

IMAGE CLASSIFICATION - A CNN APPROACH

In []: project by

Hitaishi Bairapureddy
Nikhila Reddy Vantari
Archana Gadicharla

In [1]: *# Import necessary Libraries*

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline

import tensorflow as tf
from tensorflow.keras.datasets import cifar10
from tensorflow.keras.utils import to_categorical
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import (
    Dense, Conv2D, MaxPool2D, Flatten, Dropout, BatchNormalization
)
from tensorflow.keras.callbacks import EarlyStopping
from tensorflow.keras.preprocessing.image import ImageDataGenerator

from sklearn.metrics import ConfusionMatrixDisplay, classification_report, confusion_matrix

# Load the CIFAR-10 dataset
(train_images, train_labels), (test_images, test_labels) = cifar10.load_data()

print(f"Train images shape: {train_images.shape}")
print(f"Train labels shape: {train_labels.shape}")
print(f"Test images shape: {test_images.shape}")
print(f"Test labels shape: {test_labels.shape}")

# Define class labels
class_names = ['airplane', 'automobile', 'bird', 'cat', 'deer',
               'dog', 'frog', 'horse', 'ship', 'truck']

# Visualize the training data
rows, cols = 10, 10
fig, axes = plt.subplots(rows, cols, figsize=(15, 15))
```

```

axes = axes.ravel()

num_train = len(train_images)
for i in np.arange(0, rows * cols):
    idx = np.random.randint(0, num_train)
    axes[i].imshow(train_images[idx])
    label_idx = int(train_labels[idx])
    axes[i].set_title(class_names[label_idx], fontsize=8)
    axes[i].axis('off')

plt.subplots_adjust(hspace=0.5)
plt.show()

# Plot class distribution
unique_classes, counts = np.unique(train_labels, return_counts=True)
plt.barh(class_names, counts)
plt.title('Class Distribution in Training Set')
plt.show()

unique_classes, counts = np.unique(test_labels, return_counts=True)
plt.barh(class_names, counts)
plt.title('Class Distribution in Testing Set')
plt.show()

# Data Preprocessing: Scale and one-hot encode the labels
train_images = train_images / 255.0
test_images = test_images / 255.0

train_labels_ohe = to_categorical(train_labels, 10)
test_labels_ohe = to_categorical(test_labels, 10)

# Model Building
input_shape = (32, 32, 3)
kernel_size = (3, 3)

cnn_model = Sequential()

# First Convolutional Block
cnn_model.add(Conv2D(32, kernel_size=kernel_size, activation='relu', padding='same', input_shape=input_shape))
cnn_model.add(BatchNormalization())
cnn_model.add(Conv2D(32, kernel_size=kernel_size, activation='relu', padding='same'))
cnn_model.add(BatchNormalization())
cnn_model.add(MaxPool2D(pool_size=(2, 2)))
cnn_model.add(Dropout(0.25))

```

```

# Second Convolutional Block
cnn_model.add(Conv2D(64, kernel_size=kernel_size, activation='relu', padding='same'))
cnn_model.add(BatchNormalization())
cnn_model.add(Conv2D(64, kernel_size=kernel_size, activation='relu', padding='same'))
cnn_model.add(BatchNormalization())
cnn_model.add(MaxPool2D(pool_size=(2, 2)))
cnn_model.add(Dropout(0.25))

# Third Convolutional Block
cnn_model.add(Conv2D(128, kernel_size=kernel_size, activation='relu', padding='same'))
cnn_model.add(BatchNormalization())
cnn_model.add(Conv2D(128, kernel_size=kernel_size, activation='relu', padding='same'))
cnn_model.add(BatchNormalization())
cnn_model.add(MaxPool2D(pool_size=(2, 2)))
cnn_model.add(Dropout(0.25))

# Dense Layers
cnn_model.add(Flatten())
cnn_model.add(Dense(128, activation='relu'))
cnn_model.add(Dropout(0.25))
cnn_model.add(Dense(10, activation='softmax'))

# Compile the model
metrics_list = [
    'accuracy',
    tf.keras.metrics.Precision(name='precision'),
    tf.keras.metrics.Recall(name='recall')
]

cnn_model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=metrics_list)
cnn_model.summary()

# Early Stopping Callback
early_stop = EarlyStopping(monitor='val_loss', patience=2)

# Data Augmentation
batch_size = 32
data_augmentor = ImageDataGenerator(width_shift_range=0.1, height_shift_range=0.1, horizontal_flip=True)
train_data_gen = data_augmentor.flow(train_images, train_labels_ohe, batch_size=batch_size)

steps_per_epoch = train_images.shape[0] // batch_size

# Train the model
history = cnn_model.fit(
    train_data_gen,

```

```

    epochs=50,
    steps_per_epoch=steps_per_epoch,
    validation_data=(test_images, test_labels_one),
    callbacks=[early_stop]
)

# Model Evaluation
plt.figure(figsize=(12, 16))

plt.subplot(4, 2, 1)
plt.plot(history.history['loss'], label='Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.title('Loss Evolution')
plt.legend()

plt.subplot(4, 2, 2)
plt.plot(history.history['accuracy'], label='Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.title('Accuracy Evolution')
plt.legend()

plt.subplot(4, 2, 3)
plt.plot(history.history['precision'], label='Precision')
plt.plot(history.history['val_precision'], label='Validation Precision')
plt.title('Precision Evolution')
plt.legend()

plt.subplot(4, 2, 4)
plt.plot(history.history['recall'], label='Recall')
plt.plot(history.history['val_recall'], label='Validation Recall')
plt.title('Recall Evolution')
plt.legend()

plt.show()

# Evaluate on test set
evaluation = cnn_model.evaluate(test_images, test_labels_one)
print(f'Test Accuracy: {evaluation[1] * 100:.2f}%')

# Confusion Matrix
predicted_labels = np.argmax(cnn_model.predict(test_images), axis=1)
confusion_mtx = confusion_matrix(test_labels, predicted_labels)

disp = ConfusionMatrixDisplay(confusion_matrix=confusion_mtx, display_labels=class_names)
fig, ax = plt.subplots(figsize=(10, 10))

```

```
disp.plot(xticks_rotation='vertical', ax=ax, cmap='viridis')
plt.show()
```

```
print(classification_report(test_labels, predicted_labels))
```

```
# Save the model
```

```
cnn_model.save('cnn_model_cifar10.h5')
```

Downloading data from <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz>

170498071/170498071  **6s** 0us/step

Train images shape: (50000, 32, 32, 3)

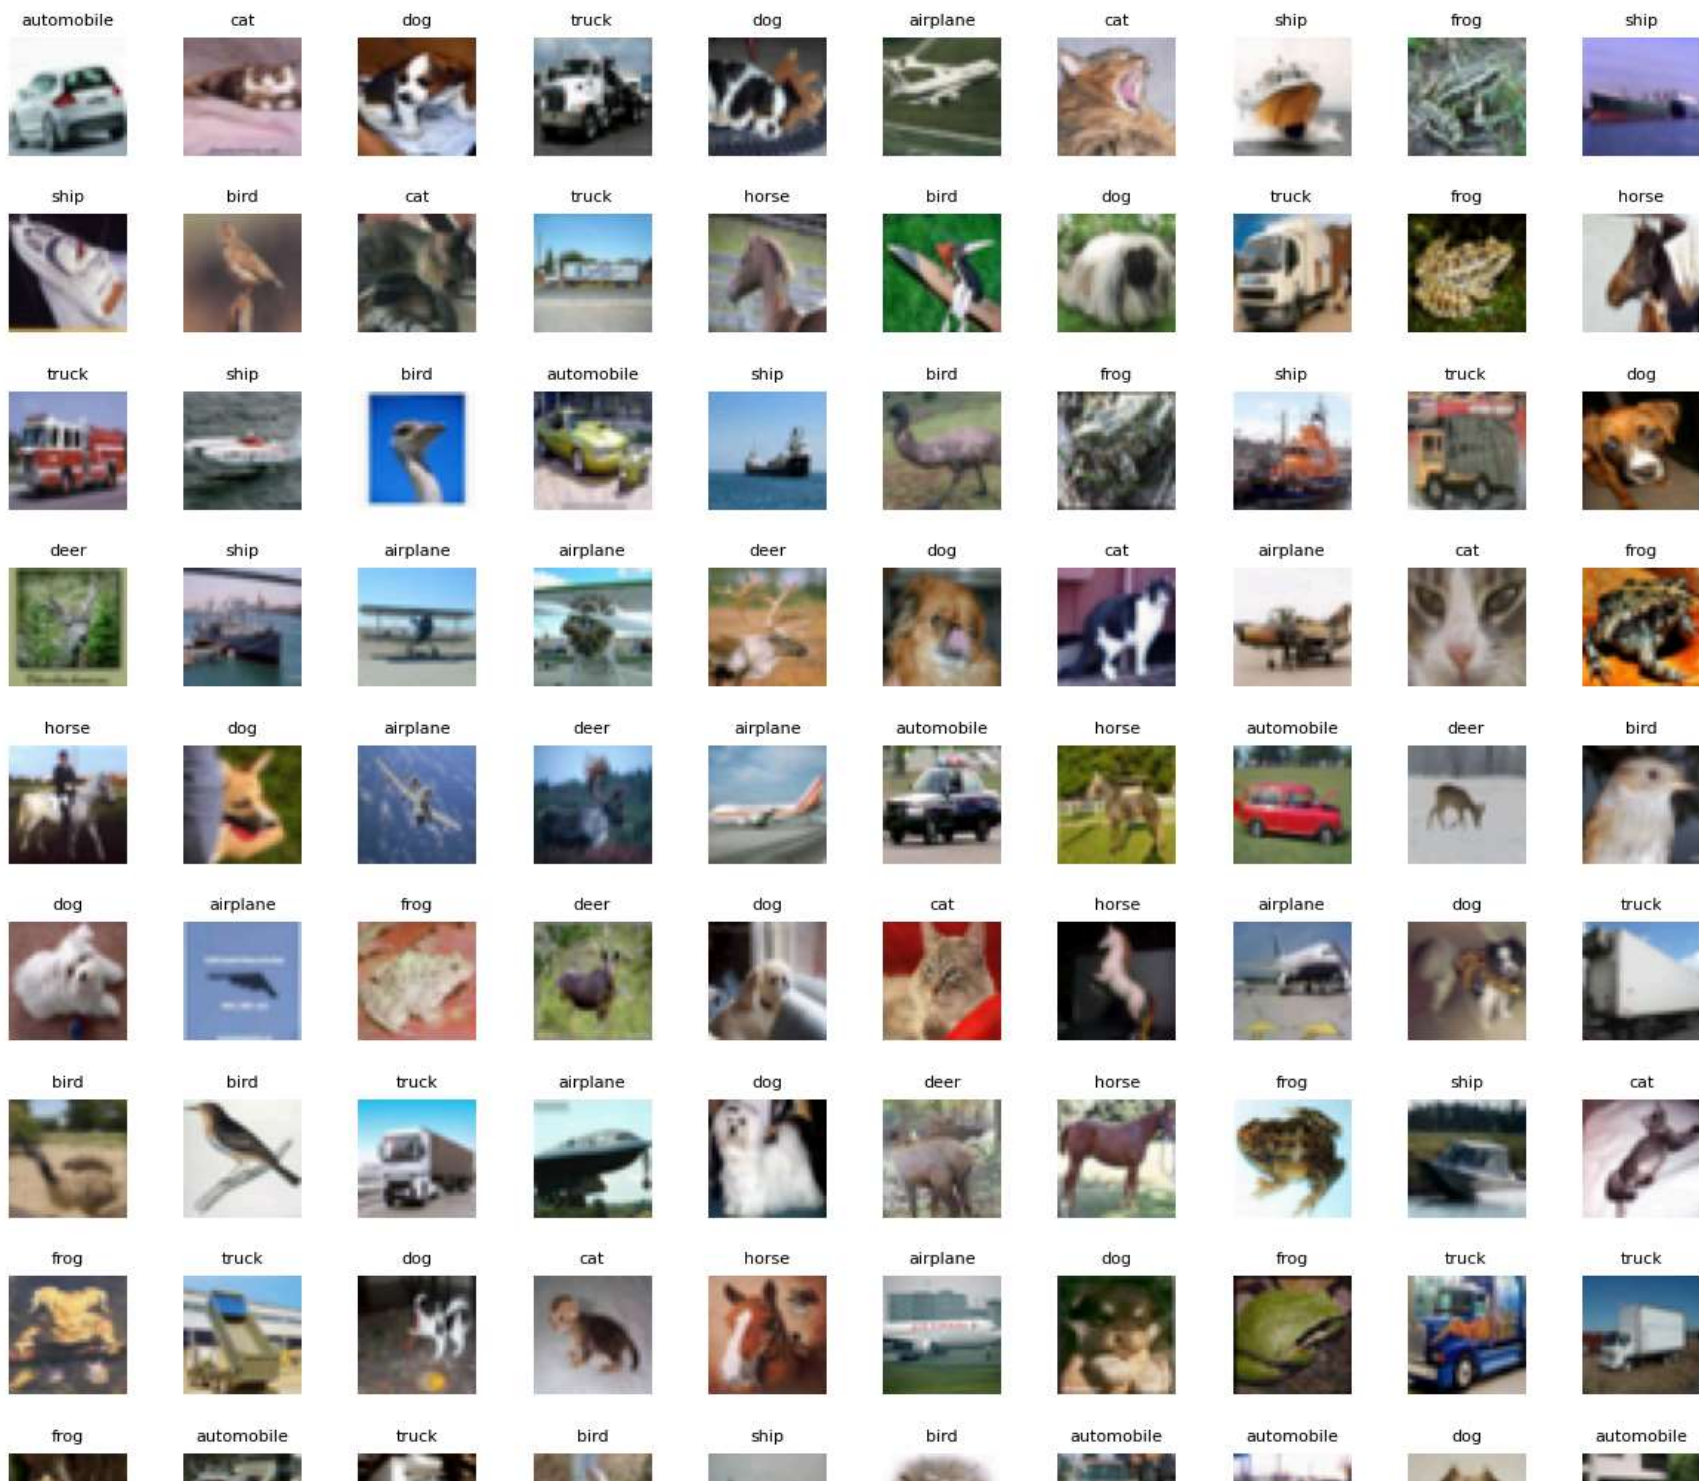
Train labels shape: (50000, 1)

Test images shape: (10000, 32, 32, 3)

Test labels shape: (10000, 1)


<ipython-input-1-c32f621ee113>:40: DeprecationWarning: Conversion of an array with ndim > 0 to a scalar is deprecated, and will error in future. Ensure you extract a single element from your array before performing this operation. (Deprecated NumPy 1.25.)

```
label_idx = int(train_labels[idx])
```





Class Distribution in Training Set





```
/usr/local/lib/python3.10/dist-packages/keras/src/layers/convolutional/base_conv.py:107: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.
```

```
    super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 32, 32, 32)	896
batch_normalization (BatchNormalization)	(None, 32, 32, 32)	128
conv2d_1 (Conv2D)	(None, 32, 32, 32)	9,248
batch_normalization_1 (BatchNormalization)	(None, 32, 32, 32)	128
max_pooling2d (MaxPooling2D)	(None, 16, 16, 32)	0
dropout (Dropout)	(None, 16, 16, 32)	0
conv2d_2 (Conv2D)	(None, 16, 16, 64)	18,496
batch_normalization_2 (BatchNormalization)	(None, 16, 16, 64)	256
conv2d_3 (Conv2D)	(None, 16, 16, 64)	36,928
batch_normalization_3 (BatchNormalization)	(None, 16, 16, 64)	256
max_pooling2d_1 (MaxPooling2D)	(None, 8, 8, 64)	0
dropout_1 (Dropout)	(None, 8, 8, 64)	0
conv2d_4 (Conv2D)	(None, 8, 8, 128)	73,856
batch_normalization_4 (BatchNormalization)	(None, 8, 8, 128)	512
conv2d_5 (Conv2D)	(None, 8, 8, 128)	147,584
batch_normalization_5 (BatchNormalization)	(None, 8, 8, 128)	512
max_pooling2d_2 (MaxPooling2D)	(None, 4, 4, 128)	0

dropout_2 (Dropout)	(None, 4, 4, 128)	0
flatten (Flatten)	(None, 2048)	0
dense (Dense)	(None, 128)	262,272
dropout_3 (Dropout)	(None, 128)	0
dense_1 (Dense)	(None, 10)	1,290

Total params: 552,362 (2.11 MB)

Trainable params: 551,466 (2.10 MB)

Non-trainable params: 896 (3.50 KB)

Epoch 1/50

```
/usr/local/lib/python3.10/dist-packages/keras/src/trainers/data_adapters/py_dataset_adapter.py:122: UserWarning: Your `PyDataset` class should call `super().__init__(**kwargs)` in its constructor. `**kwargs` can include `workers`, `use_multiprocessing`, `max_queue_size`. Do not pass these arguments to `fit()`, as they will be ignored.
  self._warn_if_super_not_called()
```


```
1562/1562 ————— 68s 36ms/step - accuracy: 0.3202 - loss: 1.8887 - precision: 0.5107 - recall: 0.1076 - val_accuracy: 0.5227 - val_loss: 1.3207 - val_precision: 0.6965 - val_recall: 0.3440
```


Epoch 2/50


```
1/1562 ————— 16s 11ms/step - accuracy: 0.5625 - loss: 1.1448 - precision: 0.6429 - recall: 0.2812
```


```
/usr/lib/python3.10/contextlib.py:153: UserWarning: Your input ran out of data; interrupting training. Make sure that your dataset or generator can generate at least `steps_per_epoch * epochs` batches. You may need to use the `.repeat()` function when building your dataset.
```


```
self.gen.throw(typ, value, traceback)
```


1562/1562  **1s** 423us/step - accuracy: 0.5625 - loss: 1.1448 - precision: 0.6429 - recall: 0.2812 - val_accuracy: 0.5209 - val_loss: 1.3370 - val_precision: 0.6968 - val_recall: 0.3390
Epoch 3/50


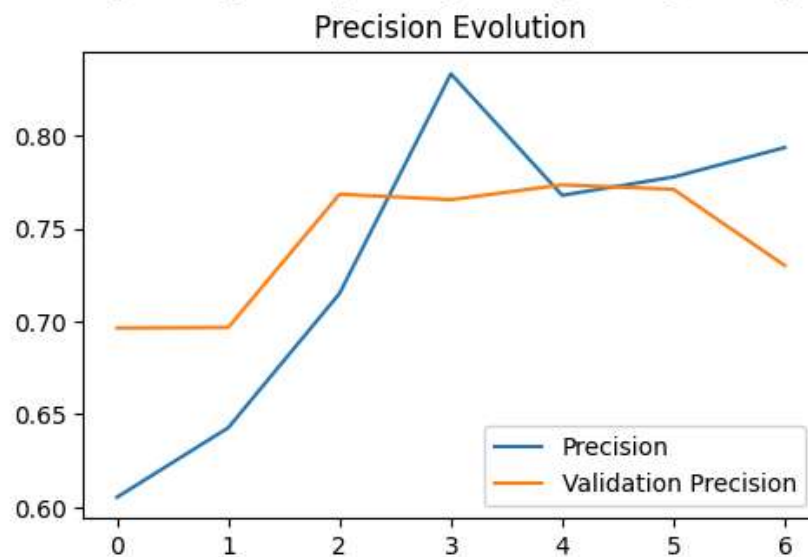

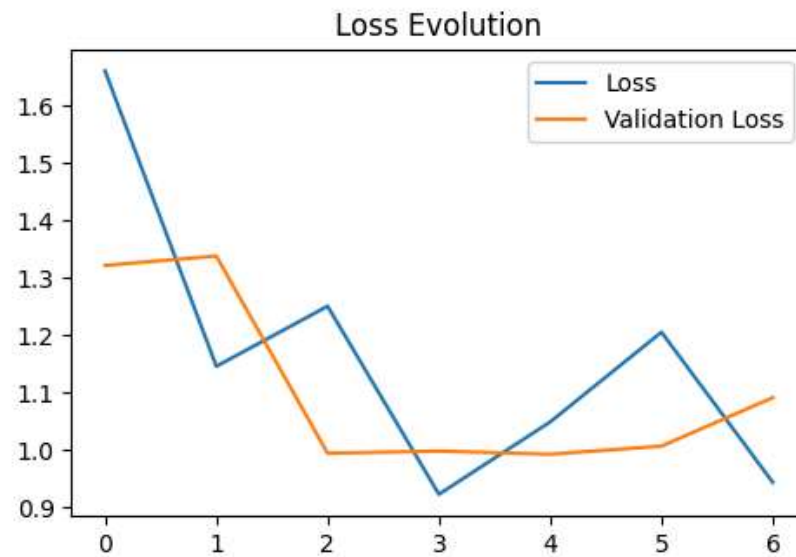
1562/1562  **65s** 30ms/step - accuracy: 0.5259 - loss: 1.3199 - precision: 0.6953 - recall: 0.3445 - val_accuracy: 0.6478 - val_loss: 0.9933 - val_precision: 0.7685 - val_recall: 0.5339
Epoch 4/50

1562/1562  **1s** 490us/step - accuracy: 0.7188 - loss: 0.9217 - precision: 0.8333 - recall: 0.4688 - val_accuracy: 0.6493 - val_loss: 0.9970 - val_precision: 0.7656 - val_recall: 0.5382
Epoch 5/50

1562/1562  **82s** 31ms/step - accuracy: 0.6181 - loss: 1.0760 - precision: 0.7604 - recall: 0.4843 - val_accuracy: 0.6724 - val_loss: 0.9915 - val_precision: 0.7736 - val_recall: 0.5917
Epoch 6/50

1562/1562  **1s** 493us/step - accuracy: 0.5625 - loss: 1.2044 - precision: 0.7778 - recall: 0.4375 - val_accuracy: 0.6681 - val_loss: 1.0056 - val_precision: 0.7712 - val_recall: 0.5867
Epoch 7/50


1562/1562  **80s** 30ms/step - accuracy: 0.6689 - loss: 0.9513 - precision: 0.7894 - recall: 0.5541 - val_accuracy: 0.6378 - val_loss: 1.0902 - val_precision: 0.7301 - val_recall: 0.5524



313/313 ————— 1s 3ms/step - accuracy: 0.6388 - loss: 1.0792 - precision: 0.7322 - recall: 0.5518

Test Accuracy: 63.78%

313/313 ————— 2s 3ms/step



airpl

automo

1

0

1

hc

1

0



Predicted label

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.

	precision	recall	f1-score	support
0	0.89	0.50	0.64	1000
1	0.90	0.72	0.80	1000
2	0.70	0.36	0.48	1000
3	0.49	0.44	0.46	1000
4	0.50	0.74	0.59	1000
5	0.78	0.34	0.48	1000
6	0.42	0.95	0.58	1000
7	0.78	0.63	0.70	1000
8	0.86	0.77	0.81	1000
9	0.69	0.91	0.78	1000
accuracy			0.64	10000
macro avg	0.70	0.64	0.63	10000
weighted avg	0.70	0.64	0.63	10000